


```
1 import os
2 import numpy as np
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5 import tensorflow as tf
6 import cv2
7 from tensorflow.keras.preprocessing.image import ImageDataGenerator
8 from tensorflow.keras.applications import MobileNetV2
9 from tensorflow.keras.layers import Dense, Flatten, Dropout, AveragePooling2D, Input
10 from tensorflow.keras.models import Model
11 from tensorflow.keras.optimizers import Adam
12 from tensorflow.keras.callbacks import EarlyStopping
13 from sklearn.metrics import confusion_matrix, classification_report
```

```
1 dataset_path = "/content/drive/MyDrive/ProgressSoft /Face Mask Detection /Face Mask Dataset" #DataSetLink:https://www.kaggle.com/datasets/ashishjangra27/face-mask-12k-images-dataset
2 train_dir = os.path.join(dataset_path, "Train")
3 val_dir = os.path.join(dataset_path, "Validation")
4 test_dir = os.path.join(dataset_path, "Test")
```

```
1 IMG_SIZE = (224, 224)
2 BATCH_SIZE = 32
3 EPOCHS = 20
4 LEARNING_RATE = 0.0001
```

```
1 #Data Augumentaion performed on training data
2 train_datagen = ImageDataGenerator(
3     rescale=1.0/255,
4     rotation_range=20,
5     zoom_range=0.15,
6     width_shift_range=0.2,
7     height_shift_range=0.2,
8     shear_range=0.15,
9     horizontal_flip=True,
10    fill_mode="nearest"
11 )
12
13 val_datagen = ImageDataGenerator(rescale=1.0/255) #Normlisation On Validation data
14 test_datagen = ImageDataGenerator(rescale=1.0/255) #Normlisation On Test data
15
```

```
1 train_generator = train_datagen.flow_from_directory(
2     train_dir, target_size=IMG_SIZE, batch_size=BATCH_SIZE, class_mode="binary"
3 )
4
5 val_generator = val_datagen.flow_from_directory(
6     val_dir, target_size=IMG_SIZE, batch_size=BATCH_SIZE, class_mode="binary"
7 )
8
9 test_generator = test_datagen.flow_from_directory(
10    test_dir, target_size=IMG_SIZE, batch_size=BATCH_SIZE, class_mode="binary", shuffle=False
11 )
12
```

 Found 10000 images belonging to 2 classes.
Found 800 images belonging to 2 classes.
Found 992 images belonging to 2 classes.

Class 0 for wearing a Mask and Class 1 for not wearing A mask

```
1 # Function to visualize some training images
2 def visualize_images(generator, title):
3     images, labels = next(generator)
4     fig, axes = plt.subplots(3, 5, figsize=(10, 6))
5     fig.suptitle(title, fontsize=16)
6     for i, ax in enumerate(axes.flat):
7         ax.imshow(images[i])
8         ax.set_title(f"Label: {int(labels[i])}")
9         ax.axis("off")
10    plt.show()
```

```
1 # Visualize Training and Validation Images
2 visualize_images(train_generator, "Sample Training Images")
3 visualize_images(val_generator, "Sample Validation Images")
```




Sample Training Images



Sample Validation Images



```
1 # load the MobileNetV2 network, ensuring the head FC layer sets are left off
2 base_model = MobileNetV2(weights="imagenet", include_top=False, input_tensor=Input(shape=(224, 224, 3)))
```

 <ipython-input-8-082d1beab1e3>:1: UserWarning: `input_shape` is undefined or non-square, or `rows` is not in [96, 128, 160, 192, 224]. Weights for input shape (224, 224) will be loaded as the default.

```
base_model = MobileNetV2(weights="imagenet", include_top=False, input_tensor=Input(shape=(224, 224, 3)))
```

```
1 # Define custom classifier on top of MobileNetV2
2 x = base_model.output
3 x = AveragePooling2D(pool_size=(7, 7))(x)
4 x = Flatten(name="flatten")(x)
5 x = Dense(128, activation="relu")(x)
6 x = Dropout(0.5)(x)
7 predictions = Dense(1, activation="sigmoid")(x) # Binary classification (Mask or No Mask)
```

```
1 # loop over all layers in the base model and freeze them so they will *not* be updated during the first training process
2 for layer in base_model.layers:
3     layer.trainable = False
```

```
1 model = Model(inputs=base_model.input, outputs=predictions)
```

```
1 model.compile(loss="binary_crossentropy", optimizer=Adam(learning_rate=LEARNING_RATE), metrics=["accuracy"])
```


✓ Saving Checkpoints

```
1 full_path = os.path.join(dataset_path, "training_1")

1 os.makedirs(full_path, exist_ok=True)

1 checkpoint_path = os.path.join(full_path, "cp.weights.h5") # Since you're only saving weights, you should use the .weights.h5 extension. If you're saving the whole model, you would use the .keras extension instead
2 checkpoint_dir = os.path.dirname(checkpoint_path)
3
4 # Create a callback that saves the model's weights
5 cp_callback = tf.keras.callbacks.ModelCheckpoint(filepath=checkpoint_path,
6                                                  save_weights_only=True,
7                                                  verbose=1)
```

✓ Training

```
1 early_stop = EarlyStopping(monitor="val_loss", patience=5, restore_best_weights=True)
2
3 history = model.fit(
4     train_generator,
5     steps_per_epoch=len(train_generator),
6     validation_data=val_generator,
7     validation_steps=len(val_generator),
8     epochs=10,
9     callbacks=[early_stop,cp_callback]
10
11 )
```

```
🔄 Epoch 1/10
313/313 [=====] - ETA: 0s - loss: 0.0985 - accuracy: 0.9659
Epoch 1: saving model to /content/drive/MyDrive/ProgressSoft /Face Mask Detection /Face Mask Dataset/training_1/cp.weights.h5
313/313 [=====] - 1911s 6s/step - loss: 0.0985 - accuracy: 0.9659 - val_loss: 0.0461 - val_accuracy: 0.9812
Epoch 2/10
313/313 [=====] - ETA: 0s - loss: 0.0718 - accuracy: 0.9756
Epoch 2: saving model to /content/drive/MyDrive/ProgressSoft /Face Mask Detection /Face Mask Dataset/training_1/cp.weights.h5
313/313 [=====] - 591s 2s/step - loss: 0.0718 - accuracy: 0.9756 - val_loss: 0.0402 - val_accuracy: 0.9850
Epoch 3/10
313/313 [=====] - ETA: 0s - loss: 0.0603 - accuracy: 0.9791
Epoch 3: saving model to /content/drive/MyDrive/ProgressSoft /Face Mask Detection /Face Mask Dataset/training_1/cp.weights.h5
313/313 [=====] - 590s 2s/step - loss: 0.0603 - accuracy: 0.9791 - val_loss: 0.0286 - val_accuracy: 0.9937
Epoch 4/10
313/313 [=====] - ETA: 0s - loss: 0.0522 - accuracy: 0.9826
Epoch 4: saving model to /content/drive/MyDrive/ProgressSoft /Face Mask Detection /Face Mask Dataset/training_1/cp.weights.h5
313/313 [=====] - 593s 2s/step - loss: 0.0522 - accuracy: 0.9826 - val_loss: 0.0266 - val_accuracy: 0.9912
Epoch 5/10
313/313 [=====] - ETA: 0s - loss: 0.0480 - accuracy: 0.9832
Epoch 5: saving model to /content/drive/MyDrive/ProgressSoft /Face Mask Detection /Face Mask Dataset/training_1/cp.weights.h5
313/313 [=====] - 594s 2s/step - loss: 0.0480 - accuracy: 0.9832 - val_loss: 0.0261 - val_accuracy: 0.9925
Epoch 6/10
313/313 [=====] - ETA: 0s - loss: 0.0419 - accuracy: 0.9865
Epoch 6: saving model to /content/drive/MyDrive/ProgressSoft /Face Mask Detection /Face Mask Dataset/training_1/cp.weights.h5
313/313 [=====] - 585s 2s/step - loss: 0.0419 - accuracy: 0.9865 - val_loss: 0.0229 - val_accuracy: 0.9912
Epoch 7/10
313/313 [=====] - ETA: 0s - loss: 0.0396 - accuracy: 0.9857
Epoch 7: saving model to /content/drive/MyDrive/ProgressSoft /Face Mask Detection /Face Mask Dataset/training_1/cp.weights.h5
313/313 [=====] - 591s 2s/step - loss: 0.0396 - accuracy: 0.9857 - val_loss: 0.0217 - val_accuracy: 0.9925
Epoch 8/10
313/313 [=====] - ETA: 0s - loss: 0.0355 - accuracy: 0.9872
Epoch 8: saving model to /content/drive/MyDrive/ProgressSoft /Face Mask Detection /Face Mask Dataset/training_1/cp.weights.h5
313/313 [=====] - 592s 2s/step - loss: 0.0355 - accuracy: 0.9872 - val_loss: 0.0200 - val_accuracy: 0.9937
Epoch 9/10
313/313 [=====] - ETA: 0s - loss: 0.0349 - accuracy: 0.9884
Epoch 9: saving model to /content/drive/MyDrive/ProgressSoft /Face Mask Detection /Face Mask Dataset/training_1/cp.weights.h5
313/313 [=====] - 591s 2s/step - loss: 0.0349 - accuracy: 0.9884 - val_loss: 0.0203 - val_accuracy: 0.9937
Epoch 10/10
313/313 [=====] - ETA: 0s - loss: 0.0327 - accuracy: 0.9882
Epoch 10: saving model to /content/drive/MyDrive/ProgressSoft /Face Mask Detection /Face Mask Dataset/training_1/cp.weights.h5
313/313 [=====] - 592s 2s/step - loss: 0.0327 - accuracy: 0.9882 - val_loss: 0.0198 - val_accuracy: 0.9925
```

```
1 os.listdir(checkpoint_dir)

🔄 ['cp.weights.h5']
```

```
1 # Loads the weights
2 model.load_weights(checkpoint_path)
```

```
1 model.compile(loss="binary_crossentropy", optimizer=Adam(), metrics=["accuracy"])
```

```
1 test_generator.reset()
2 predictions = (model.predict(test_generator) > 0.5).astype("int32")
3 true_labels = test_generator.classes

/usr/local/lib/python3.11/dist-packages/keras/src/trainers/data_adapters/py_dataset_adapter.py:121: UserWarning: Your `PyDataset` class should call `super().__init__(**kwargs)` in its constructor. `**kwargs` can include `workers`, `use_multiprocessing`, `max_queue_size`. Do not pass
  self._warn_if_super_not_called()
/usr/local/lib/python3.11/dist-packages/keras/src/models/functional.py:237: UserWarning: The structure of `inputs` doesn't match the expected structure.
Expected: ['keras_tensor']
Received: inputs=Tensor(shape=(32, 224, 224, 3))
  warnings.warn(msg)
31/31 ----- 309s 10s/step

1 print(predictions.shape, true_labels.shape)
2

(992, 1) (992, )

1 print(set(true_labels))

{np.int32(0), np.int32(1)}

1 test_loss, test_accuracy = model.evaluate(test_generator)
2 print(f"Test Accuracy: {test_accuracy*100:.2f}%")

/usr/local/lib/python3.11/dist-packages/keras/src/models/functional.py:237: UserWarning: The structure of `inputs` doesn't match the expected structure.
Expected: ['keras_tensor']
Received: inputs=Tensor(shape=(None, 224, 224, 3))
  warnings.warn(msg)
31/31 ----- 63s 2s/step - accuracy: 0.9962 - loss: 0.0192
Test Accuracy: 99.29%

1 model_path = os.path.join(dataset_path, "saved_model")
2 os.makedirs(model_path, exist_ok=True)

1 # Save the entire model as a SavedModel.
2 !mkdir -p saved_model
3 tf.saved_model.save(model, 'saved_model')

1 conf_matrix = confusion_matrix(true_labels, predictions)
2 plt.figure(figsize=(6, 5))
3 sns.heatmap(conf_matrix, annot=True, fmt="d", cmap="Blues", xticklabels=["WithMask", "WithoutMask"], yticklabels=["WithMask", "WithoutMask"])
4 plt.show()
```



```
1 # Print classification report
2 print("Classification Report:\n", classification_report(true_labels, predictions, target_names=["Mask", "No Mask"]))
3
4 # Function to visualize model predictions
5 def visualize_predictions(generator, model, title):
6     images, labels = next(generator)
7     predictions = model.predict(images)
8     predictions = np.round(predictions).astype(int).flatten()
9
10    fig, axes = plt.subplots(3, 5, figsize=(10, 6))
11    fig.suptitle(title, fontsize=16)
12
```

```
13     for i, ax in enumerate(axes.flat):
14         ax.imshow(images[i])
15         ax.set_title(f"True: {int(labels[i])}, Pred: {predictions[i]}")
16         ax.axis("off")
17
18     plt.show()
19
20 # Visualize predictions on test images
21 visualize_predictions(test_generator, model, "Test Predictions")
22
```

🔄 Classification Report:

	precision	recall	f1-score	support
Mask	0.99	1.00	0.99	483
No Mask	1.00	0.99	0.99	509
accuracy			0.99	992
macro avg	0.99	0.99	0.99	992
weighted avg	0.99	0.99	0.99	992

/usr/local/lib/python3.11/dist-packages/keras/src/models/functional.py:237: UserWarning: The structure of `inputs` doesn't match the expected structure.
Expected: ['keras_tensor']
Received: inputs=Tensor(shape=(32, 224, 224, 3))
warnings.warn(msg)

1/1 3s 3s/step

Test Predictions



```
1 def visualize_predictions(generator, model, title, samples_per_class=5):
2     # Lists to store images for each class
3     class_0_images = [] # No mask
4     class_1_images = [] # With mask
5     class_0_labels = []
6     class_1_labels = []
7
8     # Keep getting batches until we have enough samples of each class
9     while len(class_0_images) < samples_per_class or len(class_1_images) < samples_per_class:
10         images, labels = next(generator)
11         predictions = model.predict(images)
12         predictions = np.round(predictions).astype(int).flatten()
13
14         # Sort images by their true labels
15         for img, label in zip(images, labels):
16             if label == 0 and len(class_0_images) < samples_per_class: # No mask
17                 class_0_images.append(img)
18                 class_0_labels.append(label)
19             elif label == 1 and len(class_1_images) < samples_per_class: # With mask
20                 class_1_images.append(img)
21                 class_1_labels.append(label)
22
23     # Combine the collected images and labels
24     display_images = class_0_images + class_1_images
25     display_labels = class_0_labels + class_1_labels
26
27     # Make predictions on the collected images
28     predictions = model.predict(np.array(display_images))
```



```

29 predictions = np.round(predictions).astype(int).flatten()
30
31 # Plot the results
32 fig, axes = plt.subplots(2, 5, figsize=(15, 6))
33 fig.suptitle(title, fontsize=16)
34
35 for i, ax in enumerate(axes.flat):
36     ax.imshow(display_images[i])
37     ax.set_title(f"True: {' Mask' if display_labels[i]==0 else 'No Mask'}\n"
38                 f"Pred: {'Mask' if predictions[i]==0 else 'No Mask'}")
39     ax.axis("off")
40
41 plt.tight_layout()
42 plt.show()
43
44 # Visualize predictions on test images
45 visualize_predictions(test_generator, model, "Test Predictions")


```

```

1/1 ————— 2s 2s/step
1/1 ————— 2s 2s/step
1/1 ————— 2s 2s/step
1/1 ————— 2s 2s/step
1/1 ————— 2s 2s/step
1/1 ————— 2s 2s/step
1/1 ————— 2s 2s/step
1/1 ————— 2s 2s/step
1/1 ————— 2s 2s/step
1/1 ————— 2s 2s/step
1/1 ————— 2s 2s/step
1/1 ————— 2s 2s/step
1/1 ————— 2s 2s/step
1/1 ————— 2s 2s/step
1/1 ————— 2s 2s/step
1/1 ————— 2s 2s/step
1/1 ————— 2s 2s/step
1/1 ————— 2s 2s/step
/usr/local/lib/python3.11/dist-packages/keras/src/models/functional.py:237: UserWarning: The structure of `inputs` doesn't match the expected structure.
Expected: ['keras_tensor']
Received: inputs=Tensor(shape=(None, 224, 224, 3))
  warnings.warn(msg)
1/1 ————— 2s 2s/step

```


Test Predictions




True: Mask
Pred: Mask




True: Mask
Pred: Mask




True: Mask
Pred: Mask




True: Mask
Pred: Mask




True: Mask
Pred: Mask




True: No Mask
Pred: No Mask




True: No Mask
Pred: No Mask



True: No Mask
Pred: No Mask



True: No Mask
Pred: No Mask



True: No Mask
Pred: No Mask

```

1 import tensorflow as tf
2 import numpy as np
3 from tensorflow.keras.preprocessing.image import load_img, img_to_array
4 import matplotlib.pyplot as plt
5 import cv2
6
7 def test_single_image(model, image_path, img_height=224, img_width=224):
8     """
9     Enhanced test function for mask detection model with additional preprocessing
10     and more detailed output
11     """
12     # Load and preprocess the image
13     img = load_img(image_path, target_size=(img_height, img_width))
14     img_array = img_to_array(img)

```

```
15 # Additional preprocessing steps
16 # Convert to BGR (since MobileNetV2 was trained on BGR images)
17 img_array_bgr = img_array[..., ::-1]
18
19
20 # Normalize
21 img_array = img_array_bgr / 255.0
22
23 # Add batch dimension
24 img_array = np.expand_dims(img_array, axis=0)
25
26 # Make prediction
27 prediction = model.predict(img_array, verbose=0) # Disable verbose output
28 pred_class = np.round(prediction).astype(int)[0][0]
29 confidence = prediction[0][0]
30
31 # Prepare labels (0 = mask, 1 = no mask)
32 label = "No Mask" if pred_class == 1 else "Mask"
33 conf_percentage = confidence if pred_class == 1 else (1 - confidence)
34
35 # Create figure with two subplots
36 fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(12, 5))
37
38 # Original image
39 ax1.imshow(img)
40 ax1.set_title('Original Image')
41 ax1.axis('off')
42
43 # Preprocessed image
44 ax2.imshow(img_array_bgr[0] / 255.0) # Display normalized BGR image
45 ax2.set_title('Preprocessed Image\n(Model Input)')
46 ax2.axis('off')
47
48 # Add prediction info as figure suptitle
49 plt.suptitle(f'Prediction: {label} (Confidence: {conf_percentage:.2%})\n' +
50             f'Raw Model Output: {prediction[0][0]:.4f}',
51             color='green' if conf_percentage > 0.9 else 'red',
52             y=1.05)
53
54 plt.tight_layout()
55 plt.show()
56
57 # Print detailed analysis
58 print("\nDetailed Analysis:")
59 print("-" * 50)
60 print(f"Prediction Class: {label}")
61 print(f"Confidence: {conf_percentage:.2%}")
62 print(f"Raw Model Output: {prediction[0][0]:.4f}")
63 print("-" * 50)
64 print("Interpretation Guide:")
65 print("• Raw output close to 0 → Model thinks mask is present")
66 print("• Raw output close to 1 → Model thinks no mask is present")
67 if conf_percentage > 0.9:
68     print("\nNote: High confidence prediction (>90%)")
69
70 return pred_class, confidence
```

```
1 # Example usage
2 image_path = "/content/profile picture.jpg"
3 prediction, confidence = test_single_image(model, image_path)
```

Prediction: Mask (Confidence: 98.04%)
Raw Model Output: 0.0196

Original Image



Preprocessed Image
(Model Input)



Detailed Analysis:

Prediction Class: Mask
Confidence: 98.04%
Raw Model Output: 0.0196